

The earlier (1,1) point now "floats up" leading to an infinite number of planes (linear decision boundaries in 3-D) now separating the two classes (much like the concentric circles doughnut 'floating up' over the vada)

The 'Factorisation' in Math/Summation

Where does this appear, and why? **Short Answer: Everywhere!**

Working Rule:-


try putting it everywhere, and remove it if it is not required!

→ Probability (conditional/non-conditional)

→ Chain Rule (Calculus)
— total derivate/partial derivative.

PROBABILITY

The general probability factorisation



$$P(\text{cause \# } i | \text{effect}) = \frac{P(\text{effect} | \text{cause \# } i) P(\text{cause \# } i)}{P(\text{effect})}$$

Symmetrical

$$\underbrace{P(A|B) P(B)}_{P(A \text{ and } B)} = \underbrace{P(B|A) P(A)}_{P(B \text{ and } A)}$$

$$\boxed{P(A|B)} = \frac{P(B|A) \boxed{P(A)}}{P(B)}$$

a posteriori

probability of A

updated probability of A

$$x = x + 1$$

$$\Rightarrow x_{\text{new}} = x_{\text{old}} + 1$$

$$P(A)_{\text{updated}} = [\quad] \times P(A)_{\text{initial}} \rightarrow P(A)$$

$P(A|B)$

→ As such, there is no difference between a conditional probability and an unconditional probability → these are just the updated and initial variants of the same physical quantity.

$$p(A|B) = \frac{p(B|A) p(A)}{p(B)}$$

$$\boxed{p(A|B, C)} = \frac{p(B|A, C)}{p(B|C)} \boxed{p(A|C)}$$

updated prob of A initial prob of A


$$p(A)_{\text{updated}} \leftarrow \boxed{\text{Bayes Rule}} \leftarrow p(A)_{\text{initial}}$$

Where do the summations come in

$$p(\text{effect}) = \sum_j p(\text{effect} | \text{cause} \# j) p(\text{cause} \# j)$$

$\forall j$ → safely put a summation for all j 's

Example: →



Suppose we find the lawn to be wet in the morning.
 [Bayes Rule] $p(\text{rained last night} | \text{lawn wet in the morning})$

$$p(\text{rain}|\text{wet}) = \frac{p(\text{wet}|\text{rain}) p(\text{rain})}{p(\text{wet})}$$

$$= \frac{p(\text{wet}|\text{rain}) p(\text{rain})}{p(\text{wet}|\text{rain}) p(\text{rain}) + p(\text{wet}|\text{tap on}) p(\text{tap on})}$$

Impossible to visualise $I(x, y, t) \rightarrow 4\text{-D entity}$
 Use the first order Taylor series approximation.

$$f(\underline{x} + \delta \underline{x}) = f(\underline{x}) + \nabla f \cdot \delta \underline{x}$$

$$\underbrace{f(\underline{x} + \delta \underline{x}) - f(\underline{x})}_{\substack{\delta f \text{ or } \delta f(\underline{x}) \\ \text{or } \delta f(x, y, z)}} = \underbrace{\begin{bmatrix} \partial f / \partial x \\ \partial f / \partial y \\ \partial f / \partial z \end{bmatrix}}_{\substack{\nabla f \\ \text{or } \nabla f(x, y, z)}} \cdot \begin{bmatrix} \delta x \\ \delta y \\ \delta z \end{bmatrix}$$

$$= \frac{\partial f}{\partial x} \delta x + \frac{\partial f}{\partial y} \delta y + \frac{\partial f}{\partial z} \delta z$$

$$= \sum_{\text{var} = x, y, z} \left(\frac{\partial f}{\partial \text{var}} \right) (\delta \text{var})$$

Generalise to a function of D variables

$$f(\underline{x}), \quad \underline{x} = \begin{bmatrix} x_D \\ \vdots \\ x_1 \end{bmatrix} \quad \text{1st order Taylor series expansion}$$

$$\left. \begin{array}{l} \delta f \text{ or } \delta f(\underline{x}) \\ \text{or, } \delta f(x_1 \dots x_D) \end{array} \right\} = \nabla f \cdot \delta \underline{x} = \sum_{i=1}^D \frac{\partial f}{\partial x_i} \delta x_i$$

Take-home point: The total change is always a summation

$$\delta f = \delta f(\underline{x}) = \sum_{i=1}^D \frac{\partial f}{\partial x_i} \delta x_i$$

Consider another variable t

(all the x_i 's are functions of this variable t)

$$\frac{\delta f}{\delta t} = \sum_{i=1}^D \frac{\partial f}{\partial x_i} \frac{\delta x_i}{\delta t}$$

We can take the limit as $\delta t \rightarrow 0$

$$\frac{\partial f}{\partial t} = \sum_{i=1}^D \frac{\partial f}{\partial x_i} \frac{\partial x_i}{\partial t}$$

Here, we had one variable t , so the partial derivative is also the total derivative.

$$\frac{df}{dt} = \sum_{i=1}^D \frac{\partial f}{\partial x_i} \frac{dx_i}{dt}$$

Now, consider a set of variables $t_j : j \in \{1, D\}$

(all the x_i 's are functions of t_j)

$i \in \{1, D\}$

$j \in \{1, D\}$

$$\forall j : \frac{\partial f}{\partial t_j} = \sum_{i=1}^D \frac{\partial f}{\partial x_i} \frac{\partial x_i}{\partial t_j} \quad \text{Chain Rule}$$

Compact Moral of the story: if f depends on many x_i , then

$$\frac{\partial f}{\partial t} = \sum_{x_i} \frac{\partial f}{\partial x_i} \frac{\partial x_i}{\partial t}$$

Perceptron & MLP: some closing notes

- (*) Key difference: MLP and the perceptron:
MLP uses continuous sigmoidal non-linearities in the hidden units, whereas the perceptron uses a step function non-linearities
- (*) Variants: Skiplayers: either direct connection, or with a small first layer weight (so that over its operating range, the hidden unit is effectively linear), compensating with a large weight value from the hidden unit to the output
- (*) Sparse network (CNN)

Math overview/recap: → 1st order Taylor series approximation

$$\underbrace{E(\underline{w} + \delta \underline{w})}_{\substack{\text{error function} \\ \text{minimise}}} = E(\underline{w}) + \underbrace{(\nabla E) \cdot (\delta \underline{w})}_{\substack{\text{weights} \\ \text{(All)}}}$$

Overall aim: to find a weight vector, which minimises an error function $E(\underline{w})$

At the extremum, $\nabla E = \underline{0}$ (vector)

max/min/saddle point

scalar

vector

$$\begin{bmatrix} \partial E / \partial w_2 \\ \partial E / \partial w_1 \end{bmatrix}_{(2-D)} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$$

Local Quadratic Approximation

$$E(\underline{w} + \delta \underline{w}) = E(\underline{w}) + \underbrace{\nabla E \cdot (\delta \underline{w})}_{(\delta \underline{w})^T \nabla E} + \frac{1}{2} (\delta \underline{w})^T H (\delta \underline{w}) + \text{higher order terms}$$

↙ Hessian

Example: first order

$$I(x+p, y+q, t+v) = I(x, y, t) + p I_x + q I_y + v I_t$$

Second order term

$$H_{ij} = \frac{\partial^2 E}{\partial w_i \partial w_j} \Big|_{\underline{w}}$$

$$E(\underline{w} + \delta \underline{w}) = E(\underline{w}) + (\delta \underline{w})^T \nabla E + \frac{1}{2} (\delta \underline{w})^T H (\delta \underline{w})$$

= 0, at the extremum

Extremum:

$$E(\underline{w} + \delta \underline{w}) = E(\underline{w}) + \frac{1}{2} (\delta \underline{w})^T H (\delta \underline{w})$$

→ a geometric interpretation

BACK PROPAGATION

$$\underline{w}^{(t+1)} = \underline{w}^{(t)} - \eta \nabla E$$

1849, Cauchy

Error function
What is this?

Gradient Descent

We typically start with random weights

$$\begin{bmatrix} \hat{w}_{ji}^{(1)} \\ \vdots \\ \hat{w}_{ji}^{(2)} \\ \vdots \\ k_{ij} \end{bmatrix}$$

Training:
(Supervised)

[training] inputs \rightarrow outputs
To learn the weights \equiv function
which the NN implements


Error: b/w the ideal output
& the one which the NN
gives us currently.

N points
index $n = 1 \dots N$

$$E(\underline{w}) = \sum_{n=1}^N E_n(\underline{w})$$

$$\frac{\partial E(\underline{w})}{\partial \omega} = \sum_n \frac{\partial E_n}{\partial \omega}$$

Example:




Assumptions: sum of-squares errors

- Hidden layer activation function $h(a) = \tanh(a)$
- Output layer activation function $\sigma(a) = a, \forall k = a_k$

$$h(a) = \tanh(a) = \frac{e^a - e^{-a}}{e^a + e^{-a}}$$

$$\begin{aligned} h'(a) &= \frac{\partial h(a)}{\partial a} = \frac{(e^a + e^{-a})(e^a - (-e^{-a})) - (e^a - e^{-a})(e^a - e^{-a})}{(e^a + e^{-a})^2} \\ &= \frac{(e^a + e^{-a})^2 - (e^a - e^{-a})^2}{(e^a + e^{-a})^2} = 1 - \left[\frac{e^a - e^{-a}}{e^a + e^{-a}} \right]^2 = 1 - h^2(a) \end{aligned}$$

For the n 'th training data item:



$$E_n \triangleq \frac{1}{2} \sum_{k=1}^K (\gamma_k - t_k)^2$$

ground truth/
label


$\gamma_k = k$ 'th item of
the function
 $\gamma(\underline{w}, \underline{x})$

$t_k = k$ 'th item of
the target vector

why only γ_k ?
(defined only at the output
layer)

A representative example

BACKPROPAGATION (contd.) [An example]



Assume
 $\rightarrow h(\cdot)$: activation function is $\tanh(\cdot)$

$\rightarrow \sigma(\cdot)$: activation function is a unit function

[2 layers of connections]

Why is this called 'Backpropagation'?



Training phase: error is at the output: this is fed back to update the weights

$n \in \{1, N\}$

FOR each pattern in the training set in turn, we follow the following operations.

① A 'forward propagation'

(j: hidden layer)

$$a_j^{(1)} = \underline{w}_j^{(1)T} \underline{x} = \sum_{i=0}^D w_{ji}^{(1)} x_i$$

$$z_j = \tanh(a_j)$$

$$y_k = \underline{w}_k^{(2)T} \underline{z} = \sum_{j=0}^M w_{kj}^{(2)} z_j$$

② Evaluate δ_k 's for each output unit

$$\delta_k = y_k - t_k \quad \text{What is this, and how?}$$

$$E_n \triangleq \frac{1}{2} \sum_{k=1}^K (y_k - t_k)^2$$

$$\delta_k \triangleq \frac{\partial E_n}{\partial a_k} = \frac{1}{2} \cdot 2 (y_k - t_k) \left(\frac{\partial y_k}{\partial a_k} \right) = 1 \text{ as } y_k = a_k \text{ (} \sigma(\cdot) = \text{unit fn.} \text{)}$$

output layer activation

$\delta_k = y_k - t_k$ (Else, according to the specific activation function $\sigma(\cdot)$ at the output layer)

③ Backpropagate these to obtain δ_j 's for the hidden layer units

$$\delta_j = (1 - z_j^2) \sum_{k=1}^K w_{kj} \delta_k$$

What is this, and how?
[hidden layer] [output layer] previous step (step #2)

$$\delta_j \triangleq \frac{\partial E_n}{\partial a_j} = \sum_k \left(\frac{\partial E_n}{\partial a_k} \right) \left(\frac{\partial a_k}{\partial a_j} \right)$$

δ_k [step #2]

$$a_k = \underline{w}^{(2)T} \underline{z} = \sum_{j=0}^M w_{kj}^{(2)} z_j = \sum_{j=0}^M w_{kj}^{(2)} h(a_j)$$

$$\Rightarrow \frac{\partial a_k}{\partial a_j} = w_{kj}^{(2)} \frac{\partial h(a_j)}{\partial a_j} = w_{kj}^{(2)} h'(a_j) = w_{kj}^{(2)} (1 - z_j^2)$$

$$\delta_j = \sum_k \delta_k w_{kj}^{(2)} (1 - z_j^2) = (1 - z_j^2) \sum_k w_{kj}^{(2)} \delta_k$$