Home Page

Title Page

Contents

Page 13 of 30

Go Back

Full Screen

Close

Quit

Regularised Least Squares

- Why? An otherwise nice model with nice properties, but gives infinite/trivial solutions
- To control overfitting
- Start: $E_D(\mathbf{w}) \stackrel{\triangle}{=} \sum_{i=1}^{N} \{t_i - \mathbf{w}^T \boldsymbol{\phi}(\mathbf{x}_i)\}^2$ (data fidelity)
- To minimise $E_D(\mathbf{w}) + \lambda E_w(\mathbf{w})$ Fidelity, weights param λ : $E_w(\mathbf{w}) = 0$. $E_w(\mathbf{w}) = \frac{1}{2}\mathbf{w}^T\mathbf{w} = \frac{1}{2}\sum_{j=0}^{M-1} w_j^2$
- Advantage: Quadratic in $\mathbf{w}$: closed-form solution
- (ML): 'weight decay': weights $\downarrow$ 0 unless supported by the data. (Stat): 'param shrinkage'
- $E \stackrel{\triangle}{=} \sum_{i=1}^{N} \{t_i - \mathbf{w}^T \boldsymbol{\phi}(\mathbf{x}_i)\}^2 + \lambda \frac{1}{2}\mathbf{w}^T\mathbf{w}$. $\frac{\partial E}{\partial \mathbf{w}} = 0 \implies$ $\frac{-2}{2}\sum_{i=1}^{N}\{t_i - \mathbf{w}^T\boldsymbol{\phi}(\mathbf{x}_i\}\boldsymbol{\phi}^T(\mathbf{x}_i) + \frac{2\lambda}{2}\mathbf{w}^T = 0 \implies$ $\sum_{i=1}^{N} t_i \boldsymbol{\phi}^T(\mathbf{x}_i) = \sum_{i=1}^{N} \mathbf{w}^T\boldsymbol{\phi}(\mathbf{x}_i)\boldsymbol{\phi}^T(\mathbf{x}_i) + \lambda\mathbf{w}^T \implies$

- $\sum_{i=1}^{N} t_i \boldsymbol{\phi}^T(\mathbf{x}_i) = \mathbf{w}^T(\sum_{i=1}^{N} \boldsymbol{\phi}(\mathbf{x}_i)\boldsymbol{\phi}^T(\mathbf{x}_i) + \lambda \mathbf{I}) \implies$

$$= \mathbf{w}^T \; (\qquad + \lambda \mathbf{I}) \implies$$

$$\mathbf{t}^T \boldsymbol{\Phi} = (\boldsymbol{\Phi}^T\boldsymbol{\Phi} + \lambda \mathbf{I})^T \mathbf{w} = (\boldsymbol{\Phi}^T\boldsymbol{\Phi} + \lambda \mathbf{I})\mathbf{w} \implies$$

- $\boxed{\mathbf{w} = (\boldsymbol{\Phi}^T\boldsymbol{\Phi} + \lambda \mathbf{I})^{-1}\boldsymbol{\Phi}^T\mathbf{t}}$
- Note about the $||\mathbf{w}^T\mathbf{w}||$: May actually be implemented numerically as $||\mathbf{w}^T\mathbf{w}|| - c$, small c

Intertwined Histories

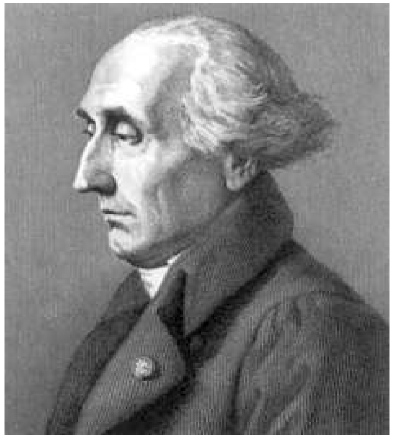

J.-L. Lagrange
[1736-1813]

A. Lavoisier
[1743-1794]

J.-B. J. Fourier
[1768-1830]

https://upload.wikimedia.org/wikipedia/commons/1/19/Lagrange_portrait.jpg

https://upload.wikimedia.org/wikipedia/commons/4/44/Lavoisier-statue.jpg

https://upload.wikimedia.org/wikipedia/commons/a/aa/Fourier2.jpg

$\mathscr{C}$lassification

- $\mathbf{x} \rightarrow$ [Classifier] $\rightarrow \mathscr{C}_j$
- Three approaches to Classification:

1. Simplest: Discriminant Functions:
   Functions which directly assign a class to $\mathbf{x}$.
   Linear Discriminant: the discriminant fns are lines/linear/hyperplanes
2. Model them directly: e.g., Mixture of Gaussians. Represent as parametric models, optimise params using a training set
3. Toughest: Generative Approach: Find $P(\mathscr{C}_j|\mathbf{x})$
   Find $P(\mathscr{C}_j|\mathbf{x})$ using the Bayes' Theorem:
   $P(\mathscr{C}_j|\mathbf{x}) = P(\mathbf{x}|\mathscr{C}_j)P(\mathscr{C}_j)/P(\mathbf{x})$. Models for:
   $P(\mathbf{x}|\mathscr{C}_j)$: class cond densities; $P(\mathscr{C}_j)$: priors

Home Page

Title Page

Contents

Page 17 of 30

Go Back

Full Screen

Close

Quit

Men of God. . .

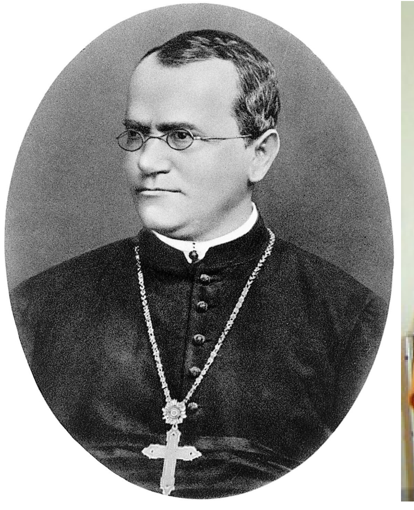

Thomas Bayes [1701-1761] G. J. Mendel [1882-1884] M. Mitra [1968-]

https://upload.wikimedia.org/wikipedia/commons/d/d4/Thomas_Bayes.gif

https://upload.wikimedia.org/wikipedia/commons/3/3d/Gregor_Mendel_oval.jpg

http://iseeindia.com/wordpress/wp-content/uploads/2011/11/Ramkrishna_Miss11736-290x290.jpg

Mahan Maharaj/Swami Vidyanathananda
2011 Shanti Swarup Bhatnagar Award in Math Sciences
Infosys Prize 2015 for Mathematical Sciences

Home Page

Title Page

Contents

Page 18 of 30

Go Back

Full Screen

Close

Quit

Useful Generalisations of Linearity

- Linearity: Written equivalently in two ways: $y(\mathbf{x},\mathbf{w}) = \mathbf{w}^T\mathbf{x}$, for $(D+1) = M-$ dim data, $x_0 = 1$, or $y(\mathbf{x},\mathbf{w}) = \mathbf{w}^T\mathbf{x} + w_0$, for $D = (M-1)-$ dim data
- $y(\mathbf{x},\mathbf{w}) = w_0x_0 + \ldots w_{M-1}x_{M-1} = \sum_{j=0}^{M-1} w_jx_j$
- Model useful for Regression: linear comb of basis fns (lin/non-lin) $\boxed{y(\mathbf{x},\mathbf{w}) = \mathbf{w}^T\boldsymbol{\phi}(\mathbf{x})}$
- Generalising lin to scalar basis functions $\phi_j(\mathbf{x})$:
- $y(\mathbf{x},\mathbf{w}) = \sum_{j=0}^{M-1} w_j\phi_j(\mathbf{x}) = w_o\phi_0(\mathbf{x}) + \ldots w_{M-1}\phi_{M-1}(\mathbf{x})$
- Model useful for Classification: fns (lin/non-lin) of the linear $\mathbf{w}^T\mathbf{x}$ (or $\mathbf{w}^T\boldsymbol{\phi}(\mathbf{x})$) $\boxed{y(\mathbf{x},\mathbf{w}) = f(\mathbf{w}^T\mathbf{x})}$
- Examples: Linear Regression, Neural Networks

Home Page
Title Page
Contents
Page 19 of 30
Go Back
Full Screen
Close
Quit

Discriminant Functions: 2 Classes

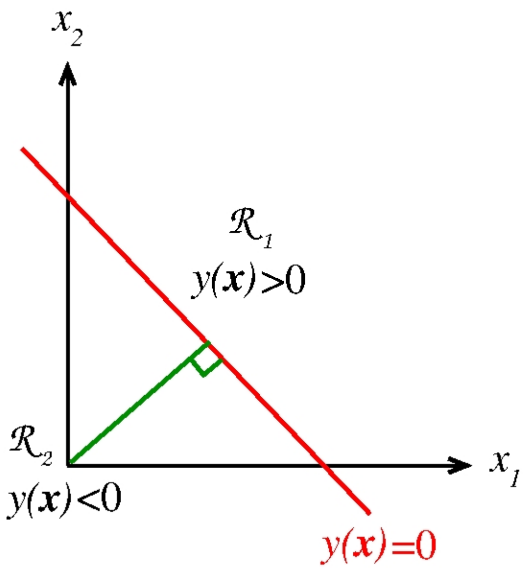

- $y(\mathbf{x}) = \mathbf{w}^T\mathbf{x} + w_0$
- 2-D implicit from of eqn of a line is $ax+by+c=0$. Here, $y(\mathbf{x},\mathbf{w}) = w_2x_2 + w_1x_1 + w_0 = 0$
- $y(\mathbf{x}) = 0$: 1-D h'plane in 2-D
- Relative location of $\mathscr{R}_1$, $\mathscr{R}_2$ is immaterial: which is above/-below/to the left/to the right
- Physical Significance of w_0: measure of the dist from the origin Why? For $ax+by+c=0$, perp distance of (x_1,y_1) from the line is $\frac{|ax_1+by_1+c|}{\sqrt{a^2+b^2}}$
  Perp dist of $y(\mathbf{x}) = 0$ from the origin $= \frac{|w_2(0)+w_1(0)+w_0|}{\sqrt{w_2^2+w_1^2}} = \frac{|w_0|}{||\mathbf{w}||}$, $||\mathbf{w}||^2 = \mathbf{w}^T\mathbf{w} = \sum w_j^2$

Some more Physical Significance

- For two points $\mathbf{x}_A$ and $\mathbf{x}_B$ on the line $y(\mathbf{x}) = 0$:
  $y(\mathbf{x}_A) = 0 \implies \mathbf{w}^T\mathbf{x}_A + w_0 = 0$
  $y(\mathbf{x}_B) = 0 \implies \mathbf{w}^T\mathbf{x}_B + w_0 = 0$
  $\implies \mathbf{w}^T(\mathbf{x}_A - \mathbf{x}_B) = 0 \implies \mathbf{w} \perp$ line $y(\mathbf{x}) = 0$
- Phy Significance of perp dist of a point from a line

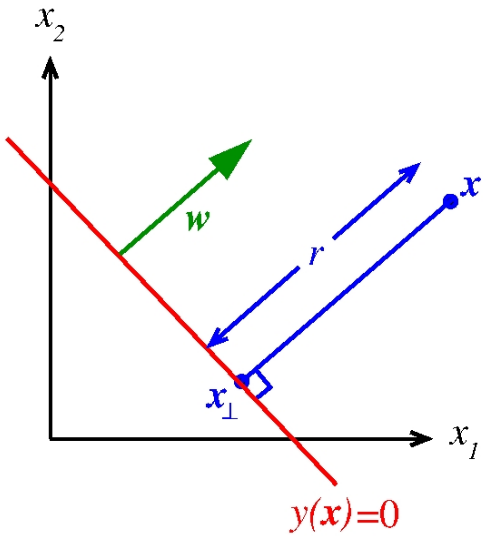

- $\mathbf{x} = \mathbf{x}_\perp + r\,\hat{\mathbf{w}} = \mathbf{x}_\perp + r\,\frac{\mathbf{w}}{||\mathbf{w}||}$
- Pre-multiply by $\mathbf{w}^T$ & add w_0:
- $\mathbf{w}^T\mathbf{x} + w_0 = \mathbf{w}^T(\mathbf{x}_\perp + r\frac{\mathbf{w}}{||\mathbf{w}||}) + w_0$
- $\implies y(\mathbf{x}) = (\mathbf{w}^T\mathbf{x}_\perp + w_0) + r\frac{||\mathbf{w}||^2}{||\mathbf{w}||}$
- $\implies \boxed{r = \frac{y(\mathbf{x})}{||\mathbf{w}||}}$
- Consistent with perp distance of (x_1, y_1) from line: $\frac{ax_1+by_1+c}{\sqrt{a^2+b^2}}$

Discriminant Functions: K Classes

- Building a $K-$ Classifier from 2-class ones

[C. M. Bishop, Pattern Recognition and Machine Learning. Springer, 2006. Fig. 4.2, p. 183]

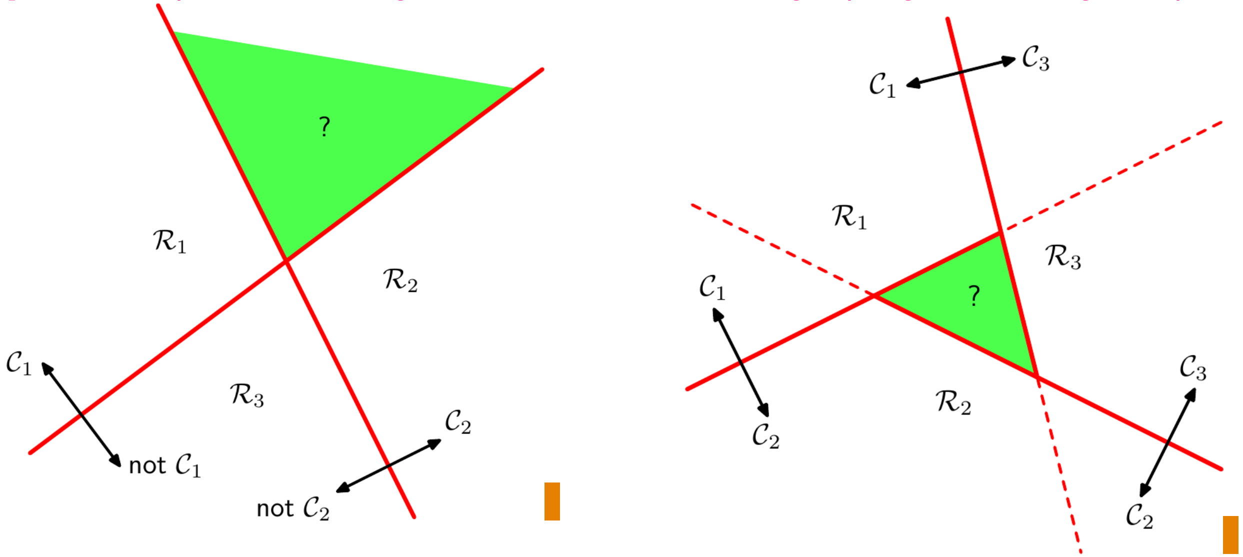

  - One-versus-Rest
  - $K-1$ classifiers, each of which solves the 2-class $\mathscr{C}_j$ vs. not $\mathscr{C}_j$

  - One-versus-One
  - ${}^K C_2$ 2-class classifiers
  - Ambiguity here also!

- e.g., Tree-SVM? Explicitly define the hierarchy!