



Home Page

Title Page

Contents

◀ ▶

◀ ▶

Page 5 of 19

Go Back

Full Screen

Close

Quit

2018 Turing Award: Deep Learning



S. Bengio

[U Montreal]

https://awards.acm.org/binaries/content/gallery/acm/awards/photo/a-b/bengio_3406375



G. Hinton

[U Toronto]

https://awards.acm.org/binaries/content/gallery/acm/awards/photo/h-j/hinton_4791679



Y. LeCun

[NYU]

https://awards.acm.org/binaries/content/gallery/acm/awards/photo/k-l/lecun_6017366

[Home Page](#)[Title Page](#)[Contents](#)[◀](#) [▶](#)[◀](#) [▶](#)[Page 6 of 19](#)[Go Back](#)[Full Screen](#)[Close](#)[Quit](#)

Perceptron Learning Criterion

- 2-class classifier, i 'th training point \mathbf{x}^i : $t_i = \pm 1$
- $y(\mathbf{x}^i) = h(\mathbf{w}^T \mathbf{x}^i) = h(a) = +1(a \geq 0), = -1(a < 0)$
- Declare Class $\mathcal{C}_1(t_i = +1)$ if $\mathbf{w}^T \mathbf{x}^i \geq 0$
- Declare Class $\mathcal{C}_2(t_i = -1)$ if $\mathbf{w}^T \mathbf{x}^i < 0$
- Combined: (SVM-like!) Correct, if $t_i \mathbf{w}^T \mathbf{x}^i \geq 0$
- Perceptron penalty: 0 if correct, else sum
- Simple criterion: $E(\mathbf{w}) \triangleq -\sum_{\forall i \in \mathcal{M}} t_i \mathbf{w}^T \mathbf{x}^i$
- no penalty for 'how much' misclassification
- Solved numerically, Cauchy's weight update rule


[Home Page](#)[Title Page](#)[Contents](#)[◀](#)[▶](#)[◀](#)[▶](#)[Page 7 of 19](#)[Go Back](#)[Full Screen](#)[Close](#)[Quit](#)

Iterative Weight Update: Learning

https://upload.wikimedia.org/wikipedia/commons/d/d3/Augustin-Louis_Cauchy_1901.jpg



- A.-L. Cauchy [1789-1857]
- **Cauchy's Rule**[1849]
- $\mathbf{w}^{(\tau+1)} = \mathbf{w}^{(\tau)} - \eta \nabla E(\mathbf{w})$
- dim cons?
- Step ' η '?
- why '-'?




- $s = \frac{E(w + \delta w) - E(w)}{(w + \delta w) - (w)} = \frac{E(w + \delta w) - E(w)}{\delta w}$
- $\lim_{\delta w \rightarrow 0} : \partial E / \partial w \implies \nabla E(\mathbf{w})$
- $\mathbf{w}: [w_j], \nabla E(\mathbf{w}) = [\partial E / \partial w_j]$
- '-' sign: go against the gradient!

- ' η ': step size or learning rate, tuning (adaptive?)
- small η : small steps, longer attain local min
- large η : large steps, may miss local min

[Home Page](#)[Title Page](#)[Contents](#)[Page 8 of 19](#)[Go Back](#)[Full Screen](#)[Close](#)[Quit](#)






Multi-Layer Perceptron



[Home Page](#)[Title Page](#)[Contents](#)[Page 9 of 19](#)[Go Back](#)[Full Screen](#)[Close](#)[Quit](#)

Activation Functions

- Neuron input: scalar, sum of weighted inputs
- Activation fn: possible non-linearity, scalar output
- X'fer fns: sigmoid, tanh, ReLU, Leaky ReLU, eLU



- $LeakyRelU(a, \alpha) \triangleq \max(\alpha a, a), \alpha \in (0, 1)$

- $eLU(a, \alpha) \triangleq \begin{cases} a, & a > 0 \\ \alpha(e^a - 1), & a \leq 0 \end{cases}$ [00:50]